

# Introduction to Regression Analysis

Jerome Dumortier

17 August 2023

## Empirical research approach

- Scientific method

## Types of regression models used in social sciences

- Explain the mechanics of the linear regression model
- Assumptions necessary to obtain unbiased estimates from the classical linear regression model
- Number of applications for this type of model are infinite.

## Review of statistical concepts

- Population versus sample
- Hypothesis testing

Introduction  
to Regression  
Analysis

Jerome  
Dumortier

Empirical  
Research  
Approach

Types of  
Regression  
Models

Review of  
Statistical  
Concepts

# Empirical Research Approach

# Hypothesis, Model, and Data

The empirical research approach in social sciences consists of multiple steps:

① Statement of theory or hypothesis:

- People increase their consumption when their income increases by less than the income increase.

② Specification of the mathematical model

$$\Delta consumption = \beta_0 + \beta_1 \cdot \Delta income$$

③ Obtaining the data

- Real personal consumption expenditures per capita
- Real Disposable Personal Income: Per Capita

④ Estimation of the parameters of the econometric model

# Estimation Results

```
##  
## Call:  
## lm(formula = diff(consumption) ~ diff(income), data = cindata)  
##  
## Residuals:  
##      Min       1Q   Median       3Q      Max   
## -745.80 -121.34  -19.25   178.05   595.55   
##  
## Coefficients:  
##              Estimate Std. Error t value Pr(>|t|)      
## (Intercept)  196.2986    77.7556   2.525   0.016 *      
## diff(income)   0.5648     0.1021   5.534 2.68e-06 ***  
## ---  
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1  
##  
## Residual standard error: 296.6 on 37 degrees of freedom  
## Multiple R-squared:  0.4529, Adjusted R-squared:  0.4381  
## F-statistic: 30.63 on 1 and 37 DF,  p-value: 2.683e-06
```

# Post-Estimation Procedures

## ⑥ Hypothesis testing

- Are the results aligned with the theory.

## ⑦ Forecasting or prediction

- How will a predicted increase in income affect consumption?

## ⑧ Using the model for control or policy purposes.

- What are the effects of stimulus spending on the economy?

The aforementioned approach is at the core of social science research such as public affairs, criminal justice, economics, or sociology. Examples:

- What are the determinants of tax avoidance with respect to cigarette consumption?
- What is the relationship between automatic bill payment and electricity consumption?

# Types of Regression Models

# Classical Linear Regression Model (CLRM)

Bivariate regression model:

- One dependent variable (e.g., home value) and one independent variable (e.g., square footage)

$$homevalue_i = \beta_0 + \beta_1 \cdot sqft_i + \epsilon_i$$

- Useful to explain the mechanics of ordinary least square models

Multivariate regression model

- One dependent variable (e.g., home value) and multiple independent variables (e.g., square footage, bedrooms)

$$homevalue_i = \beta_0 + \beta_1 \cdot sqft_i + \beta_2 \cdot bedrooms_i + \epsilon_i$$



## CLRM: Assumptions

Required assumptions for unbiasedness coefficient estimates:

- A1: Linear regression model, i.e., linear in terms of coefficients
- A2: Zero mean value of error terms  $\epsilon$ , i.e.,  $E(\epsilon_i|x_i) = 0$
- A3: Homoscedasticity (equal variance) of  $\epsilon_i$ , i.e.,  $Var(\epsilon_i) = \sigma^2$
- A4: No autocorrelation between the error terms, i.e.,  $Cov(\epsilon_i, \epsilon_j) = 0$
- A5: No covariance between  $\epsilon_i$  and  $x_i$
- A6: Number of observations is greater than number of parameters to be estimated
- A7: Full rank and absence of (perfect) multicollinearity

## CLRM: Relaxing Assumptions

Relaxing the assumptions of the classical regression model requires regression diagnostics and/or different regression approaches

- Multicollinearity: Correlation between independent variables are correlated with each other?
  - Beds and bathrooms in a home value model
  - Multicollinearity occurs between two or more (!) independent variables
- Heteroscedasticity: Errors variance not constant
- Autocorrelation between error terms
- Inclusion of irrelevant or exclusion of relevant independent variables

# Qualitative Choice Models

## Binary choice model (Probit and Logit models)

- Dependent variable: 1 or 0 (“yes” or “no”)
- Example: Recidivism (i.e., committing a crime after release from prison)

## Ordered logit

- More than two categories of the dependent variable but ordered
- Example: Level of support for a particular policy, e.g., strongly opposed, opposed, neutral, supportive, strongly supportive.

## Multinomial logit

- More than two categories of the dependent variable but no order
- Example: Commute to campus by bike, bus, car, or on foot

# Panel Data and Time Series

## Panel data regression models

- Observation of the same individual or unit over multiple years
- Fixed effects versus random effects models

## Dynamic models

- Autoregressive and distributed-lag models

## Time series model

## More Models

### Limited dependent variable models

- Some problem with the dependent variable.
- Truncation versus censoring
- Count regression models
- Duration/Hazard/Survival models

### Simultaneous equations models

- More than one dependent variable
- System of equations

### Instrumental variable models

- Omitted variables
- Measurement errors
- Simultaneous causality

# Review of Statistical Concepts

# Population versus Sample

## Population versus sample

- The population is characterized by parameters that will always remain unknown.
- Given a sample taken from the population allows us to learn something about the population parameters.
- The sample needs to be drawn at random.

The sample mean is the arithmetic average of the values in a random sample. It is usually denoted

$$\bar{x} = \frac{x_1 + x_2 + \cdots + x_n}{n} = \frac{1}{n} \sum_{i=1}^n x_i$$

## Measures of Dispersion

Population variance

$$\sigma^2 = \frac{\sum_{i=1}^N (x_i - \mu)^2}{N}$$

Sample variance

$$s^2 = \frac{\sum_{i=1}^N (x_i - \bar{x})^2}{N - 1}$$

Standard deviation is the square root of the variance, i.e.,  $\sigma = \sqrt{\sigma^2}$  or  $s = \sqrt{s^2}$ .



# Sampling Variance

The sampling variance is expressed as:

$$\text{Var}(\bar{X}_n) = \frac{\sigma^2}{n}$$

This is different from the sample variance! The sampling variance represents the variation of a particular statistic, e.g., mean. The larger  $n$ , the smaller the variance in the mean of the various drawings.

# Confidence Intervals

## Definition

- A 95% confidence interval for a parameter is an interval obtained from a sample that has a 95% probability of producing a interval containing the true value of the parameter.

## Computation:

$$\bar{x} \pm t_{df, \alpha} \cdot \frac{s}{\sqrt{n}}$$

Example data of starting salary (in 1,000) after college graduation:

Student	1	2	3	4	5	6	7	8	9	10
Salary	87	43	59	64	59	71	73	49	68	65

## Calculation of a Confidence Interval

Given the data salary, we have  $\bar{x} = 63.8$ ,  $s = 12.44$ , and  $t_{9,0.025} = 2.262$ . Using the equation:

$$\bar{x} \pm t_{df,\alpha} \cdot \frac{s}{\sqrt{n}}$$

and plugging in the data:

$$63.8 \pm 2.262 \cdot \frac{12.44}{\sqrt{10}} = 63.8 \pm 8.90$$

## Confidence Interval with R

```
salary = c(87,43,59,64,59,71,73,49,68,65)  
t.test(salary)
```

```
##  
## One Sample t-test  
##  
## data: salary  
## t = 16.225, df = 9, p-value = 5.695e-08  
## alternative hypothesis: true mean is not equal to 0  
## 95 percent confidence interval:  
## 54.90474 72.69526  
## sample estimates:  
## mean of x  
## 63.8
```

## Hypothesis Testing

A hypothesis is a statement about a parameter taking on a specific value. A hypothesis test is a procedure to verify the statement and the steps are:

- ① Formulating the null hypothesis  $H_0$  stating that the parameter takes a specific value:
  - One-sided test:  $H_0: \mu \geq \mu_0$  or  $\mu \leq \mu_0$
  - Two-sided test:  $H_0: \mu = \mu_0$
- ② Setting the significance level  $\alpha$ , e.g., 1%, 5%, or 10%.
- ③ Test statistic: Value based on the sample used to **reject** or **fail to reject** the null hypothesis.
- ④ Critical value and  $p$ -value:
  - Critical value represents the threshold between rejecting and failing to reject  $H_0$ .
  - $p$ -Value: Probability of observing the parameter given the null hypothesis. Small  $p$ -values represent evidence against  $H_0$ .

Note that equality is always part of  $H_0$ , i.e.,  $=$ ,  $\leq$ , or  $\geq$ .

## Decisions and Errors in Hypothesis Testing

Null Hypothesis	Fail to reject $H_0$	Reject $H_0$
$H_0$ is true	Correct	Type I Error
$H_0$ is false	Type II Error	Correct

Type I Error:

- Probability of rejecting  $H_0$  when it is true.
- Also known as the significance level of a test denoted with  $\alpha$ .

Type II Error:

- Probability of failing to reject  $H_0$  when it is false.

## Interpretation of the $p$ -Value

Each statistical software provides a  $p$ -value:

- Lowest level of significance at which the null hypothesis can be rejected.
- Represents the probability of observing the sample given that the hypothesis is true. The lower the  $p$ -value the more unlikely is the hypothesis.
- The null hypothesis  $H_0$  is rejected if the  $p$ -value is smaller than the significance level.

The smaller the  $p$ -value, the stronger the evidence against  $H_0$  being true. This is true for any type of hypothesis test.

## Two-sided and One-sided Hypothesis Tests

### Two-sided test

- $H_0: \mu = \mu_0$  and  $H_a: \mu \neq \mu_0$
- Reject  $H_0$  if  $|t| > t_{\alpha/2, n-1}$

### One-sided test (left-sided)

- $H_0: \mu \leq \mu_0$  and  $H_a: \mu > \mu_0$
- Reject  $H_0$  if  $|t| > t_{\alpha, n-1}$

### One-sided test (right-sided)

- $H_0: \mu \geq \mu_0$  and  $H_a: \mu < \mu_0$
- Reject  $H_0$  if  $|t| > t_{\alpha, n-1}$

In both cases,  $|t|$  refers to the absolute value of the test statistic. Two-sided tests are of importance in regression analysis.



# Hypothesis Testing: One-sample vs. Two-sample

## One-sample (or one-group) tests

- Population proportion
- Population mean with unknown variance

## Two-sample (or two-group) tests

- Population proportions
- Population means (differentiation between equal and unequal variance)
- Paired difference test

Note: Textbooks often include “population mean with *known* variance.” This is a highly unlikely case and thus, it is skipped.

# One-Group Proportion

Test statistic:

$$z = \frac{\bar{p} - p_0}{\sqrt{p_0 \cdot (1 - p_0)/n}}$$

where  $p_0$  is the hypothesized population proportion.

# One-Group Proportion in R

```
##  
## One Sample t-test  
##  
## data: gss$owngun  
## t = -1.8174, df = 1888, p-value = 0.06932  
## alternative hypothesis: true mean is not equal to 0.3333333  
## 95 percent confidence interval:  
## 0.2929756 0.3348698  
## sample estimates:  
## mean of x  
## 0.3139227
```

# One-Group Mean

Unknown variance requires the use of the  $t$ -distribution given the following test statistic:

$$t = \frac{\bar{x} - \mu}{s/\sqrt{n}}$$

where  $\bar{x}$  is the sample mean,  $\mu$  is the hypothesized mean,  $s$  is the sample standard deviation, and  $n$  is the sample size.

## One-Group Mean in R

```
##  
## One Sample t-test  
##  
## data: eggweights$weight  
## t = 1.8378, df = 60, p-value = 0.07104  
## alternative hypothesis: true mean is not equal to 60  
## 95 percent confidence interval:  
## 59.90723 62.19113  
## sample estimates:  
## mean of x  
## 61.04918
```

## Two-Group Proportions

Hypothesis test for difference between two population proportions

$$H_0 : p_1 - p_2 = 0$$

```
t.test(gss$owngun~gss$female)
```

```
##  
##  Welch Two Sample t-test  
##  
## data:  gss$owngun by gss$female  
## t = 3.8992, df = 1731.9, p-value = 0.0001002  
## alternative hypothesis: true difference in means between group 0 and group 1 is not equal to 0  
## 95 percent confidence interval:  
##  0.04184529 0.12654762  
## sample estimates:  
## mean in group 0 mean in group 1  
##      0.3608124      0.2766160
```

## Two-Group Means

Difference between two mean:

$$H_0 : \bar{x}_1 - \bar{x}_2 = 0$$

Means of two dependent populations

- Assumption of equal variance, i.e.,  $\sigma_1^2 = \sigma_2^2$
- Example: Pre- and post-test
- Pooled-Variance t-test: One estimate of unknown  $\sigma^2$ , i.e.,  $s_p$ .

Means of two independent populations

- Assumption of unequal variance, i.e.,  $\sigma_1^2 \neq \sigma_2^2$
- Samples from two different populations
- Separate-Variance t-test: Two estimates for unknown  $\sigma_1^2$  and  $\sigma_2^2$ .

## Two-Group Means: Equal Variance

```
t.test(indyhomes$price~indyhomes$zip,var.equal=TRUE)
```

```
##  
##  Two Sample t-test  
##  
## data:  indyhomes$price by indyhomes$zip  
## t = 2.0005, df = 100, p-value = 0.04816  
## alternative hypothesis: true difference in means between group 46228 and group 46268  
## 95 percent confidence interval:  
##      1510.38 363678.01  
## sample estimates:  
## mean in group 46228 mean in group 46268  
##           381600.2           199006.0
```



## Two-Group Means: Unequal Variance

```
t.test(indyhomes$price~indyhomes$zip,var.equal=FALSE)
```

```
##  
##  Welch Two Sample t-test  
##  
## data:  indyhomes$price by indyhomes$zip  
## t = 2.0403, df = 51.323, p-value = 0.04648  
## alternative hypothesis: true difference in means between group 46228 and group 46268  
## 95 percent confidence interval:  
##      2953.984 362234.402  
## sample estimates:  
## mean in group 46228 mean in group 46268  
##           381600.2           199006.0
```

## Paired Difference Test I

Difference between paired (!) values:

$$D_i = x_{1,i} - x_{2,i}$$

Elimination of variation among subjects. Point estimate for paired difference

$$\bar{D} = \frac{1}{n} \sum_{i=1}^n D_i$$

Sample standard deviation

$$S_d = \sqrt{\frac{\sum_{i=1}^n (D_i - \bar{D})^2}{n - 1}}$$

## Paired Difference Test II

Test statistic

$$t_p = \frac{\bar{D} - \mu_D}{S_d / \sqrt{n}}$$

Confidence interval

$$\bar{D} \pm t_{\alpha/2} \frac{S_D}{\sqrt{n}}$$

$t_p$  has  $n-1$  degrees of freedom

## Textbook Example

Book	Online	Bookstore	Difference
History 1	10.2	11.4	-1.2
History 2	18.95	19	-0.05
Economics 1	184.53	200.75	-16.22
Business 1	236.75	247.2	-10.45
Business 2	67.41	71.25	-3.48

Note that  $\sum D_i = -31.76$ ,  $\bar{D} = -6.352$ , and  $s_D = 6.833$ .

## Textbook Example in R

```
online      = c(10.20,18.95,184.53,236.75,67.41)
bookstore   = c(11.40,19,200.75,247.20,71.25)
t.test(online,bookstore,paired=TRUE)
```

```
##
## Paired t-test
##
## data:  online and bookstore
## t = -2.0788, df = 4, p-value = 0.1062
## alternative hypothesis: true mean difference is not equal to 0
## 95 percent confidence interval:
##  -14.835834  2.131834
## sample estimates:
## mean difference
##          -6.352
```